

MUHAMMAD GHULAM JILLANI

Lead AI & Cloud Data Scientist | Generative AI · LLMops · Agentic RAG · Multi-Cloud (AWS, Azure, GCP)

[LinkedIn](#) | [Hire Me \(Upwork\)](#) | [GitHub](#) | [Kaggle](#) | [Medium](#)
[Portfolio](#) | [Website](#) | [Book a Call](#) | m.g.jillani123@gmail.com

SELECTED RESULTS

70% faster enterprise document search · 90% ESG compliance accuracy · 80%+ of tier 1 calls automated · 48% better inventory planning · 60% faster compliance review · 93% prediction accuracy

PROFILE

I design and ship production grade Generative AI that automates complex workflows and delivers measurable ROI. Across 6+ years and 22+ end to end AI builds on AWS, Azure, and GCP, I turn fragmented data and manual operations into reliable, secure, autonomous systems for healthcare, finance, legal, retail, and SaaS. I take systems from concept to production with the evals, guardrails, and observability that keep them dependable.

WHY ENTERPRISES WORK WITH ME

- Outcomes, not demos: 22+ systems in production, with up to 75% retrieval gains, 65% less manual work, and 40% higher efficiency.
- Enterprise grade by default: PII redaction, role-based access, evals, guardrails, and monitoring built in, across GDPR, HIPAA, and SOC 2 aware workflows.
- Fits your stack: integrates with CRMs, ERPs, data warehouses, and internal tools on AWS, Azure, and GCP, including VPC and private deployment.
- Owned to production: I take systems from architecture through deployment and clean handoff, then keep them reliable.

Industries: Healthcare, finance, legal, retail, SaaS, and other regulated, reliability critical sectors.

Engagement: Discovery and architecture > Build and integrate > Evaluate and harden > Deploy, monitor, and support.

Recognition: LinkedIn Top Voice in AI and Cloud ML · Top 100 Global Kaggle Master contributor · Kaggle-X BIPOC mentor.

TECHNICAL STACK

LLMs & Reasoning: GPT-5.5 · Claude Opus 4.8 / Sonnet 4.7 · Gemini 3.1 Pro / 3.5 Flash · Llama 4 · DeepSeek R1 · Qwen 3 · reasoning and tool-use models

Agents & Orchestration: LangGraph · OpenAI Agents SDK · Google ADK · CrewAI · AutoGen · PydanticAI · LlamaIndex · n8n · MCP (Model Context Protocol)

Retrieval & Voice: Hybrid search · rerankers · GraphRAG · ColBERT · Pinecone · Weaviate · Qdrant · pgvector · LiveKit · Pipecat · Deepgram · ElevenLabs · OpenAI Realtime

LLMops & MLOps: LangSmith · Langfuse · Ragas · Arize Phoenix · evals · guardrails · vLLM · LoRA / QLoRA · MLflow · CI/CD · Docker · Kubernetes

Backend & Cloud: Python · FastAPI · Next.js · AWS Bedrock · Azure AI Foundry · GCP Vertex AI · Snowflake · Kafka · Airflow

FEATURED PROJECTS (PROBLEM · SOLUTION · RESULT)

Agentic AI & RAG Platforms

Enterprise Knowledge Navigator *LangGraph · Pinecone · GPT-4o*

Problem: Employees lost hours hunting across scattered policies, SOPs, and technical docs.

Solution: Agentic RAG where planner and retriever agents reason over a metadata tagged vector index, with role based access so each team sees only what it should.

Result: 70% less internal search time and higher productivity across support and engineering.

ESG Compliance Analyzer *LangChain · Azure OpenAI · ChromaDB*

Problem: Manual ESG and EU CSRD auditing was slow, costly, and error prone.

Solution: Agentic RAG plus LLMops that ingests disclosures, maps them to CSRD clauses, and flags gaps with citations, wrapped in monitoring for auditable runs.

Result: 90% compliance classification accuracy and 70% less manual auditing time.

FinServe GPT *CrewAI · LangGraph · AutoGen · pgvector*

Problem: Analysts could not track transactions, portfolio risk, and market signals fast enough.

Solution: Multi agent system where specialized agents watch transactions, score portfolio risk, and turn live market signals into ranked, explainable recommendations.

Result: 60% lower analyst workload with real time, AI backed insights.

AI Blog Generation Platform *LangGraph · GPT-4.1 · Tavily*

Problem: Content teams were bottlenecked writing long, research heavy technical articles.

Solution: Parallel LangGraph orchestration that fans a topic out to research, outline, and writer agents running concurrently, producing cited, multi section drafts with diagrams.

Result: 5x faster production and 80% less manual writing.

Agentic Automation & Voice AI

VoiceFlow AI *LiveKit · Deepgram · ElevenLabs · LangGraph*

Problem: Support and sales lines needed instant, human like responses around the clock.

Solution: Real time voice agent that streams speech to text, grounds every answer in your knowledge base, and triggers live CRM actions, all under one second of latency.

Result: Automated 80%+ of tier 1 calls without handoff and cut handle time 45%.

Multi-Agent Deep Research System *LangGraph · CrewAI · Tavily*

Problem: Teams spent hours compiling research across the web and internal sources.

Solution: Planner, researcher, and writer agents that decompose a question, search the web and internal data in parallel, then synthesize a sourced, structured report.

Result: Multi hour research compressed into minutes with traceable citations.

n8n + MCP Automation Suite *n8n · MCP · LangChain · FastAPI*

Problem: AI value stalled because agents were not wired into real business tools.

Solution: n8n workflows and the Model Context Protocol give LLM agents safe, governed access to CRMs, email, databases, and SaaS, so automations act, not just chat.

Result: Removes recurring manual hours across support, sales, and operations.

LLM Evaluation & Observability Platform *LangSmith · Langfuse · Ragas*

Problem: Production LLM apps regressed silently on quality and cost.

Solution: Evaluation harness and tracing layer that scores each release against golden datasets, runs regression and guardrail checks, and surfaces quality and cost on live dashboards.

Result: Catches quality and cost regressions before users ever see them.

Production Rescue Engine *Python · FastAPI · CI/CD · evals*

Problem: Companies were stuck with broken or AI generated codebases that would not ship.

Solution: A structured rescue: audit the codebase, add tests and evals, instrument observability, fix the breakages, and hand back a documented, production ready system.

Result: Failing AI builds restored to dependable, production grade systems.

AI Copilots & Conversational AI

RetailGPT Copilot *AWS Bedrock · RLHF · Snowflake*

Problem: Retailers needed personalized recommendations and smarter inventory planning.

Solution: GenAI copilot that blends real time user behavior with demand forecasts and is tuned with reinforcement learning to personalize recommendations per shopper.

Result: 32% higher conversion and 48% better inventory planning accuracy.

MedIQ Insights *Claude 4 · Vertex AI · Weaviate*

Problem: Clinicians lacked fast, grounded decision support at the point of care.

Solution: Hybrid agentic RAG that combines patient history, lab results, and current medical literature, mapped to ICD-10, to surface grounded, cited decision support.

Result: 25% better diagnostic support accuracy and 40% faster onboarding.

AI Compliance Copilot *GPT-4o (Bedrock) · OpenSearch · PII redaction*

Problem: Legal teams drowned in GDPR, HIPAA, and SOC 2 review.

Solution: Copilot that parses contracts and regulatory updates, validates them against GDPR, HIPAA, and SOC 2, and raises AI risk flags with a full audit trail.

Result: 45% lower risk exposure and 60% faster review.

Conversational Support AI *GPT-4o · AWS Bedrock · LangChain*

Problem: Customers waited too long for accurate support answers.

Solution: Context aware assistant with intent recognition and retrieval grounded answers that personalizes replies and hands off cleanly to a human when confidence drops.

Result: 40% faster query resolution and higher satisfaction.

AI Meeting Assistant (AIMA) *GPT-4o · Pinecone · Google APIs*

Problem: Teams lost time on manual notes and follow ups.

Solution: Agent that transcribes calls, summarizes decisions, and generates minutes, wired to Google Sheets for roadmaps and Gmail for automatic dispatch.

Result: 80% less manual meeting documentation.

Document Intelligence

LegalDoc AI *FAISS · GPT-4o (Bedrock) · OpenSearch*

Problem: Reviewing contracts and case files clause by clause was painfully slow.

Solution: Agentic RAG over multi format contracts and case files, with a chatbot that answers clause level questions and cites the exact source passage.

Result: 70% less review time, with clauses surfaced in seconds.

DocuWiz AI *BART · DistilBERT · RAG*

Problem: Mixed documents needed extraction, summary, and sentiment at scale.

Solution: Document intelligence pipeline that extracts, summarizes, and scores sentiment across legal, academic, and corporate files in one automated pass.

Result: 60% less review time and stronger automation.

Document Summarization Engine *Pegasus · Vertex AI*

Problem: Long PDFs and DOCX slowed legal, finance, and research work.

Solution: Transformer based summarizer that turns long PDFs and DOCX into concise, faithful summaries, preserving the key clauses reviewers care about.

Result: 60% faster document processing.

Computer Vision & Real-Time AI

Real-Time Traffic Monitoring *YOLOv7 · OpenCV · Edge AI*

Problem: Authorities needed live congestion and road safety analytics.

Solution: Computer vision pipeline that detects and counts vehicles in live video at the edge, feeding congestion and safety analytics straight to traffic operators.

Result: Better traffic flow optimization and safety insight.

Object Tracking for Security *YOLO · DeepSORT · CI/CD*

Problem: Surveillance needed reliable multi object tracking at scale.

Solution: Multi object tracking that follows people and vehicles across frames for crowd and perimeter monitoring, shipped through automated CI/CD.

Result: Automated threat detection and higher efficiency.

Plant Disease Detection *ResNet50 · Azure AI Studio*

Problem: Farmers needed early, accessible crop disease diagnosis.

Solution: Image classifier trained in the cloud and served on the web, letting a farmer photograph a crop and get an instant diagnosis plus treatment guidance.

Result: 70% faster detection and 30% better yield prediction.

Predictive ML & MLOps Platforms

AutoML-Studio *MLflow · SHAP · Cloud Run*

Problem: Non-technical teams could not build or ship ML models.

Solution: No code AutoML platform that lets non-technical users train, explain, and deploy models, with experiment tracking, SHAP explainability, and drift detection built in.

Result: Democratized model building and deployment.

AIOps Predictive Maintenance *SageMaker · Kafka · LSTM*

Problem: Unplanned equipment failures caused costly downtime.

Solution: Streaming platform that ingests sensor data through Kafka and runs real time LSTM anomaly detection to predict equipment failures before they happen.

Result: 40% less downtime and stronger maintenance efficiency.

Sales Forecasting Engine *LSTM · ARIMA · SageMaker*

Problem: Businesses lacked reliable demand and sales forecasts.

Solution: Hybrid LSTM and ARIMA forecasting served through a live dashboard, giving planners reliable demand signals for inventory and growth decisions.

Result: 30% better forecasting accuracy.

Customer Satisfaction Prediction *ZenML · MLflow · Scikit-learn*

Problem: Companies could not anticipate customer dissatisfaction.

Solution: Feedback prediction model wrapped in a reproducible ZenML and MLflow pipeline, so retraining and deployment stay one click and fully tracked.

Result: 93% prediction accuracy.

LET'S BUILD SOMETHING THAT SHIPS

Planning an AI product, agent, or automation workflow, or need a stalled build rescued into production? Let's talk.

Book a call: calendly.com/jillanisoftech/30mins **Hire me on Upwork:** [Upwork profile](#)

Website: jillanisoftech.com **Email:** m.g.jillani@jillanisoftech.com